3. Inside of the repeated 10-fold CV I ran all of my tuning for:
    1. **KNN** where k = (1:40)
    2. **Random forest** where mtry = (1:7), nodesize=(1,3,5,6,7,10), ntree=(500,1000,1500)
    3. **Neural net** where decay = (0, 0.01, 0.1, 1), size = (1, 3, 6, 10)
    4. **SVM** where cost = (5, 10), gamma = (0.1, 1)

    To compare these models, I used the validation misclassification error rate on each fold and then plotted the matrix of resultant values on a box and whiskers plot. I evaluated the best model to be the one with the lowest mean misclassification error rate and also considered the 25^{th} and 75^{th} percentiles to ensure they were also at a minimum. That is, if two models had a very similar means, I chose the one with the lower percentiles but slightly wider spread, not the one with very consistent values.